

# neural\_network

March 6, 2024

## 1 Assignment 2

In this assignment you will create a coordinate-based multilayer perceptron in numpy from scratch. For each input image coordinate  $(x, y)$ , the model predicts the associated color  $(r, g, b)$ .

You will then compare the following input feature mappings  $\gamma(\mathbf{v})$ .

- No mapping:  $\gamma(\mathbf{v}) = \mathbf{v}$ .
- Basic mapping:  $\gamma(\mathbf{v}) = [\cos(2\pi\mathbf{v}), \sin(2\pi\mathbf{v})]^T$ .
- Gaussian Fourier feature mapping:  $\gamma(\mathbf{v}) = [\cos(2\pi\mathbf{B}\mathbf{v}), \sin(2\pi\mathbf{B}\mathbf{v})]^T$ , where each entry in  $\mathbf{B} \in \mathbb{R}^{m \times d}$  is sampled from  $\mathcal{N}(0, \sigma^2)$ .

Some notes to help you with that:

- You will implement the mappings in the helper functions `get_B_dict` and `input_mapping`.
- The basic mapping can be considered a case where  $\mathbf{B} \in \mathbb{R}^{2 \times 2}$  is the identity matrix.
- For this assignment,  $d$  is 2 because the input coordinates in two dimensions.
- You can experiment with  $m$ , like  $m = 256$ .
- You should show results for  $\sigma$  value of 1.

Source: <https://bmild.github.io/fourfeat/> This assignment is inspired by and built off of the authors' demo.

### 1.1 Setup

#### 1.1.1 (Optional) Colab Setup

If you aren't using Colab, you can delete the following code cell. Replace the path below with the path in your Google Drive to the uploaded assignment folder. Mounting to Google Drive will allow you access the other .py files in the assignment folder and save outputs to this folder

```
[ ]: # you will be prompted with a window asking to grant permissions
# click connect to google drive, choose your account, and click allow
from google.colab import drive
drive.mount("/content/drive")
```

```
[ ]: # TODO: fill in the path in your Google Drive in the string below
# Note: do not escape slashes or spaces in the path string
import os
datadir = "/content/assignment2"
```

```

if not os.path.exists(datadir):
    !ln -s "/content/drive/MyDrive/CS444/assignment2-2" $datadir
os.chdir(datadir)
!pwd

```

### 1.1.2 Imports

```

[3]: import matplotlib.pyplot as plt
from tqdm.notebook import tqdm
import os, imageio
import cv2
import numpy as np

import sys
sys.path.append('/content/drive/MyDrive/CS444/assignment2-2/models')
from models.neural_net import NeuralNetwork

# makes sure your NeuralNetwork updates as you make changes to the .py file
%load_ext autoreload
%autoreload 2

# sets default size of plots
%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0)

```

## 1.2 Helper Functions

### 1.2.1 Image Data and Feature Mappings (Fill in TODOs)

```

[4]: # Data loader - already done for you
def get_image(size=512, \
              image_url='https://bmild.github.io/fourfeat/img/lion_orig.png'):

    # Download image, take a square crop from the center
    img = imageio.imread(image_url)[..., :3] / 255.
    c = [img.shape[0]//2, img.shape[1]//2]
    r = 256
    img = img[c[0]-r:c[0]+r, c[1]-r:c[1]+r]

    if size != 512:
        img = cv2.resize(img, (size, size))

    plt.imshow(img)
    plt.show()

    # Create input pixel coordinates in the unit square
    coords = np.linspace(0, 1, img.shape[0], endpoint=False)

```

```

x_test = np.stack(np.meshgrid(coords, coords), -1)
test_data = [x_test, img]
train_data = [x_test[::2, ::2], img[::2, ::2]]

return train_data, test_data

```

```

[5]: # Create the mappings dictionary of matrix B - you will implement this
def get_B_dict(size):
    mapping_size = size // 2 # you may tweak this hyperparameter
    B_dict = {}
    B_dict['none'] = None
    # add B matrix for basic, gauss_1.0
    # TODO implement this
    B_dict['basic'] = np.eye(2)
    B_gauss = np.random.normal(size=(mapping_size, 2))
    for scale in [1., 10., 100.]:
        B_dict[f'gauss_{scale}'] = B_gauss * scale

    return B_dict

```

```

[6]: # Given tensor x of input coordinates, map it using B - you will implement
def input_mapping(x, B):
    if B is None:
        # "none" mapping - just returns the original input coordinates
        return x
    else:
        # "basic" mapping and "gauss_X" mappings project input features using B
        # TODO implement this
        x_proj = (2.*np.pi*x) @ B.T
        return np.concatenate([np.sin(x_proj), np.cos(x_proj)], axis=-1)

```

### 1.2.2 MSE Loss and PSNR Error (Fill in TODOs)

```

[7]: def mse(y, p):
    # TODO implement this
    # make sure it is consistent with your implementation in neural_net.py
    return np.mean(np.square(y - p))

def psnr(y, p):
    # TODO implement this
    mse_value = mse(y, p)
    psnr_value = 10 * np.log10(255**2 / mse_value)
    return psnr_value

```

```

[8]: size = 32
train_data, test_data = get_image(size)

```

```
C:\Users\liang\AppData\Local\Temp\ipykernel_22048\3999599749.py:6:
DeprecationWarning: Starting with ImageIO v3 the behavior of this function will
switch to that of iio.v3.imread. To keep the current behavior (and make this
warning disappear) use `import imageio.v2 as imageio` or call
`imageio.v2.imread` directly.
img = imageio.imread(image_url)[..., :3] / 255.
```



Some suggested hyperparameter choices to help you start - hidden layer count: 4 - hidden layer size: 256 - number of epochs: 1000 - learning rate: 1e-4

```
[9]: # TODO: Set the hyperparameters
num_layers = 4
hidden_size = 256
hidden_sizes = [hidden_size] * (num_layers - 1)
```

```

epochs = 1000
learning_rate = 1e-4
output_size = train_data[1].shape[2]
B_dict = get_B_dict(size)

print('B_dict items:')
for k,v in B_dict.items():
    print('\t',k,np.array(v).shape)

```

```

B_dict items:
    none ()
    basic (2, 2)
    gauss_1.0 (16, 2)
    gauss_10.0 (16, 2)
    gauss_100.0 (16, 2)

```

```

[10]: # Apply the input feature mapping to the train and test data - already done for
      ↪you
def get_input_features(B_dict, mapping):
    # mapping is the key to the B_dict, which has the value of B
    # B is then used with the function `input_mapping` to map x
    y_train = train_data[1].reshape(-1, output_size)
    y_test = test_data[1].reshape(-1, output_size)
    X_train = input_mapping(train_data[0].reshape(-1, 2), B_dict[mapping])
    X_test = input_mapping(test_data[0].reshape(-1, 2), B_dict[mapping])
    return X_train, y_train, X_test, y_test

```

### 1.2.3 Plotting and video helper functions (you don't need to change anything here)

```

[11]: def plot_training_curves(train_loss, train_psnr, test_psnr):
    # plot the training loss
    plt.subplot(2, 1, 1)
    plt.plot(train_loss)
    plt.title('MSE history')
    plt.xlabel('Iteration')
    plt.ylabel('MSE Loss')

    # plot the training and testing psnr
    plt.subplot(2, 1, 2)
    plt.plot(train_psnr, label='train')
    plt.plot(test_psnr, label='test')
    plt.title('PSNR history')
    plt.xlabel('Iteration')
    plt.ylabel('PSNR')
    plt.legend()

    plt.tight_layout()

```

```

plt.show()

def plot_reconstruction(p, y):
    p_im = p.reshape(size,size,3)
    y_im = y.reshape(size,size,3)

    plt.figure(figsize=(12,6))

    # plot the reconstruction of the image
    plt.subplot(1,2,1), plt.imshow(p_im), plt.title("reconstruction")

    # plot the ground truth image
    plt.subplot(1,2,2), plt.imshow(y_im), plt.title("ground truth")

    print("Final Test MSE", mse(y, p))
    print("Final Test psnr",psnr(y, p))

def plot_reconstruction_progress(predicted_images, y, N=8):
    total = len(predicted_images)
    step = total // N
    plt.figure(figsize=(24, 4))

    # plot the progress of reconstructions
    for i, j in enumerate(range(0,total, step)):
        plt.subplot(1, N, i+1)
        plt.imshow(predicted_images[j].reshape(size,size,3))
        plt.axis("off")
        plt.title(f"iter {j}")

    # plot ground truth image
    plt.subplot(1, N+1, N+1)
    plt.imshow(y.reshape(size,size,3))
    plt.title('GT')
    plt.axis("off")
    plt.show()

def plot_feature_mapping_comparison(outputs, gt):
    # plot reconstruction images for each mapping
    plt.figure(figsize=(24, 4))
    N = len(outputs)
    for i, k in enumerate(outputs):
        plt.subplot(1, N+1, i+1)
        plt.imshow(outputs[k]['pred_imgs'][-1].reshape(size, size, -1))
        plt.title(k)
    plt.subplot(1, N+1, N+1)
    plt.imshow(gt)
    plt.title('GT')

```

```

plt.show()

# plot train/test error curves for each mapping
iters = len(outputs[k]['train_psnrs'])
plt.figure(figsize=(16, 6))
plt.subplot(121)
for i, k in enumerate(outputs):
    plt.plot(range(iters), outputs[k]['train_psnrs'], label=k)
plt.title('Train error')
plt.ylabel('PSNR')
plt.xlabel('Training iter')
plt.legend()
plt.subplot(122)
for i, k in enumerate(outputs):
    plt.plot(range(iters), outputs[k]['test_psnrs'], label=k)
plt.title('Test error')
plt.ylabel('PSNR')
plt.xlabel('Training iter')
plt.legend()
plt.show()

# Save out video
def create_and_visualize_video(outputs, size=size, epochs=epochs,
    filename='training_convergence.mp4'):
    all_preds = np.concatenate([outputs[n]['pred_imgs'].
    reshape(epochs,size,size,3)[:25] for n in outputs], axis=-2)
    data8 = (255*np.clip(all_preds, 0, 1)).astype(np.uint8)
    f = os.path.join(filename)
    imageio.mimwrite(f, data8, fps=20)

# Display video inline
from IPython.display import HTML
from base64 import b64encode
mp4 = open(f, 'rb').read()
data_url = "data:video/mp4;base64," + b64encode(mp4).decode()

N = len(outputs)
if N == 1:
    return HTML(f'''
    <video width=256 controls autoplay loop>
        <source src="{data_url}" type="video/mp4">
    </video>
    ''')
else:
    return HTML(f'''
    <video width=1000 controls autoplay loop>
        <source src="{data_url}" type="video/mp4">
    ''')

```

```

</video>
<table width="1000" cellspacing="0" cellpadding="0">
  <tr>{''.join(N*[f'<td width="{1000//len(outputs)}"></td>'])}</tr>
  <tr>{''.join(N*[f'<td style="text-align:center">{</td>'])}</tr>
</table>
''.format(*list(outputs.keys()))

```

#### 1.2.4 Experiment Runner (Fill in TODOs)

```

[12]: def NN_experiment(X_train, y_train, X_test, y_test, input_size, num_layers,\
                        hidden_size, hidden_sizes, output_size, epochs,\
                        learning_rate, opt):

    # Initialize a new neural network model
    net = NeuralNetwork(input_size, hidden_sizes, output_size, num_layers, opt)

    # Variables to store performance for each epoch
    train_loss = np.zeros(epochs)
    train_psnr = np.zeros(epochs)
    test_psnr = np.zeros(epochs)
    predicted_images = np.zeros((epochs, y_test.shape[0], y_test.shape[1]))

    # For each epoch...
    for epoch in tqdm(range(epochs)):

        # Shuffle the dataset
        # TODO implement this
        train_data = np.hstack((X_train, y_train))
        np.random.shuffle(train_data)
        X_train_shuffle = train_data[:, 0:X_train.shape[1]]
        y_train_shuffle = train_data[:, -3:]

        # Training
        # Run the forward pass of the model to get a prediction and record the
        ↪psnr
        output = net.forward(X_train_shuffle)
        train_psnr[epoch] = psnr(y_train_shuffle, output)

        # Run the backward pass of the model to compute the loss, record the
        ↪loss, and update the weights
        loss = net.backward(y_train_shuffle)
        train_loss[epoch] = loss
        net.update(lr=learning_rate, opt=opt)

        # Testing
        # No need to run the backward pass here, just run the forward pass to
        ↪compute and record the psnr

```



```

predict = net.forward(X_test)
test_psnr[epoch] = psnr(y_test, predict)
predicted_images[epoch] = predict

return net, train_psnr, test_psnr, train_loss, predicted_images

```

## 2 Low Resolution Reconstruction

### Low Resolution Reconstruction - SGD - None Mapping

```

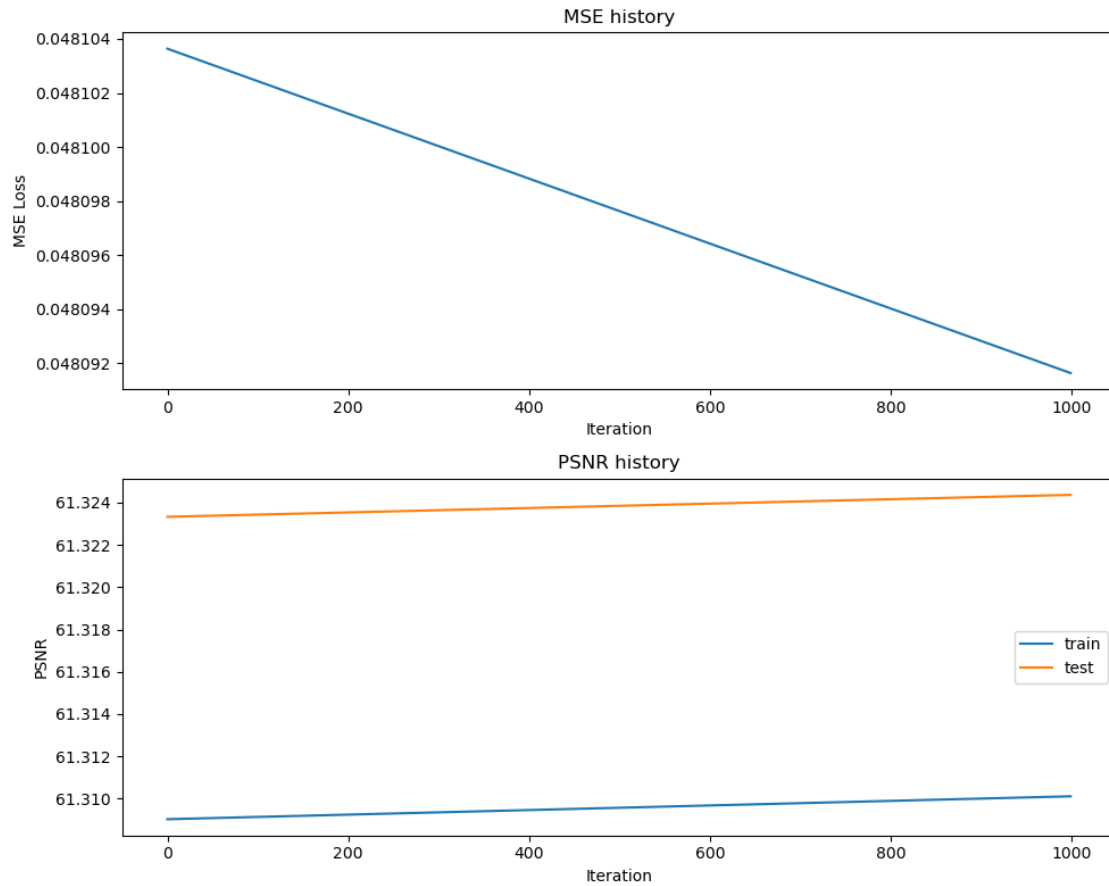
[13]: # get input features
      # TODO implement this by using the get_B_dict() and get_input_features() helper_
      ↪ functions

mapping = 'none'
B_dict = get_B_dict(128)
X_train, y_train, X_test, y_test = get_input_features(B_dict, mapping)
input_size = X_train.shape[1]
opt = "SGD"

# run NN experiment on input features
# TODO implement by using the NN_experiment() helper function
net, train_psnr, test_psnr, train_loss, predicted_images = ↪
    ↪ NN_experiment(X_train, y_train, X_test, y_test, input_size, num_layers,\
                    hidden_size, hidden_sizes, output_size, epochs,\
                    learning_rate, opt)
# plot results of experiment
plot_training_curves(train_loss, train_psnr, test_psnr)
plot_reconstruction(net.forward(X_test), y_test)
plot_reconstruction_progress(predicted_images, y_test)

0%|          | 0/1000 [00:00<?, ?it/s]

```



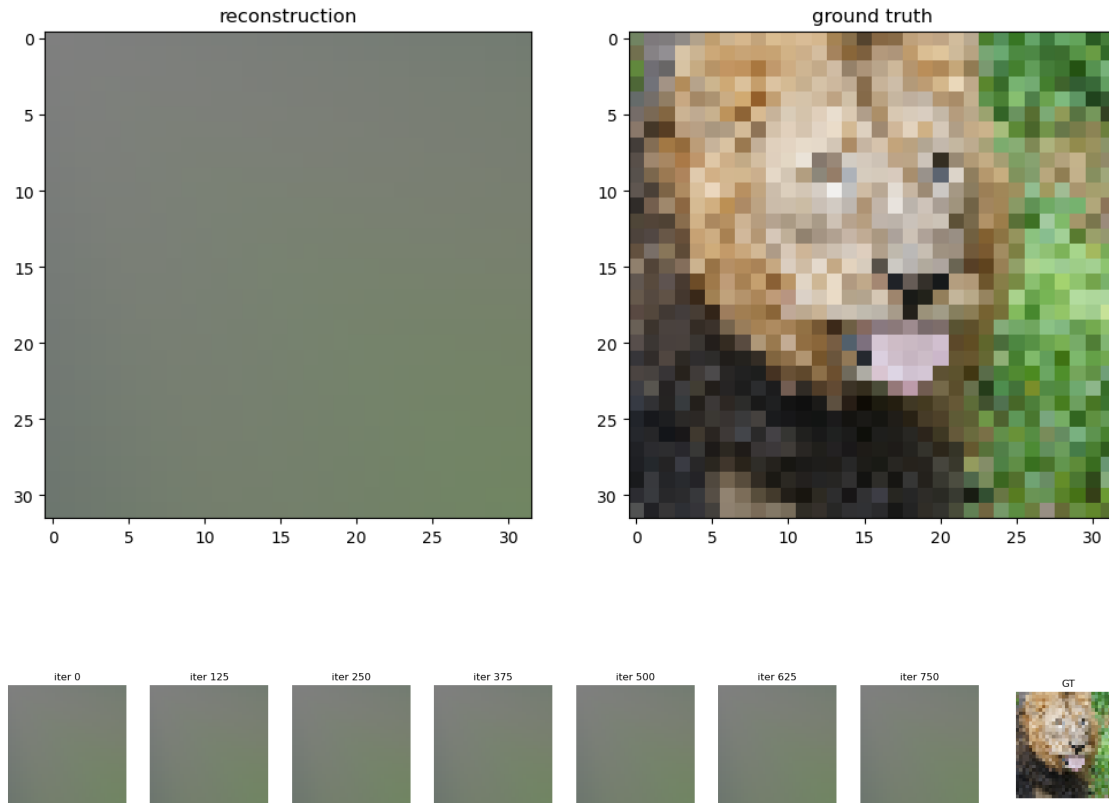
Final Test MSE 0.04793399888877405

Final Test psnr 61.324366993936586

C:\Users\liang\AppData\Local\Temp\ipykernel\_22048\3758130134.py:49:

MatplotlibDeprecationWarning: Auto-removal of overlapping axes is deprecated since 3.6 and will be removed two minor releases later; explicitly call `ax.remove()` as needed.

```
plt.subplot(1, N+1, N+1)
```




### Low Resolution Reconstruction - Adam - None Mapping

```
[14]: # get input features
# TODO implement this by using the get_B_dict() and get_input_features() helper
# functions

mapping = 'none' # None Mapping
X_train, y_train, X_test, y_test = get_input_features(B_dict, mapping)
input_size = X_train.shape[1]
opt = "Adam"
# run NN experiment on input features
# TODO implement by using the NN_experiment() helper function
net, train_psnr, test_psnr, train_loss, predicted_images = \
    NN_experiment(X_train, y_train, X_test, y_test, input_size, num_layers, \
                  hidden_size, hidden_sizes, output_size, epochs, \
                  learning_rate, opt)
# plot results of experiment
plot_training_curves(train_loss, train_psnr, test_psnr)
plot_reconstruction(net.forward(X_test), y_test)
plot_reconstruction_progress(predicted_images, y_test)
```

0%| | 0/1000 [00:00<?, ?it/s]



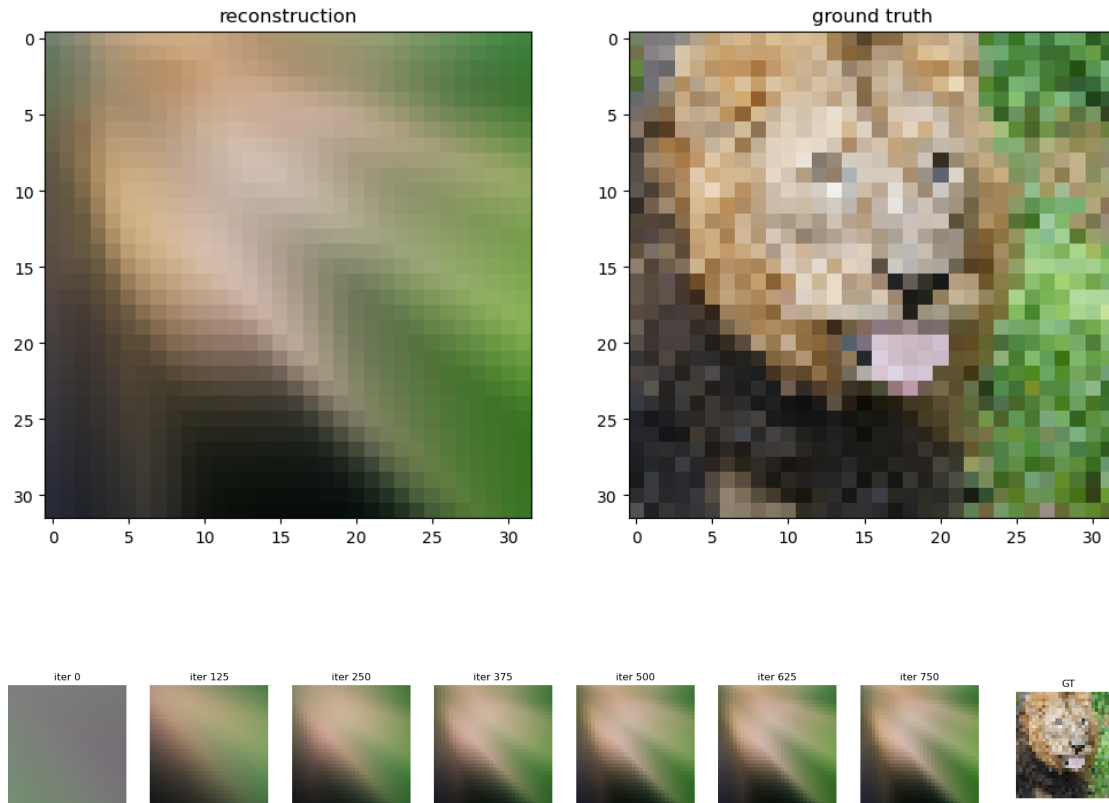
Final Test MSE 0.01515893947431834

Final Test psnr 66.32411541930865

C:\Users\liang\AppData\Local\Temp\ipykernel\_22048\3758130134.py:49:

MatplotlibDeprecationWarning: Auto-removal of overlapping axes is deprecated since 3.6 and will be removed two minor releases later; explicitly call `ax.remove()` as needed.

```
plt.subplot(1, N+1, N+1)
```



## Low Resolution Reconstruction - Optimizer of your Choice - Various Input Mapping Strategies

```
[15]: def train_wrapper(mapping, size, opt):
    # TODO implement
    # makes it easy to run all your mapping experiments in a for loop
    # this will similar to what you did previously in the last two sections

    X_train, y_train, X_test, y_test = get_input_features(B_dict, mapping)
    input_size = X_train.shape[1]
    net, train_psnrs, test_psnrs, train_loss, predicted_images = _
    ↪ NN_experiment(X_train, y_train, X_test, y_test, input_size, num_layers, \
                   hidden_size, hidden_sizes, output_size, epochs, \
                   learning_rate, opt)

    return {
        'net': net,
        'train_psnrs': train_psnrs,
        'test_psnrs': test_psnrs,
        'train_loss': train_loss,
        'pred_imgs': predicted_images
    }
```

```
[16]: outputs = {}
      for k in tqdm(B_dict):
          print("training", k)
          outputs[k] = train_wrapper(k, size, opt)
```

```
0%|          | 0/5 [00:00<?, ?it/s]
```

```
training none
```

```
0%|          | 0/1000 [00:00<?, ?it/s]
```

```
training basic
```

```
0%|          | 0/1000 [00:00<?, ?it/s]
```

```
training gauss_1.0
```

```
0%|          | 0/1000 [00:00<?, ?it/s]
```




```
training gauss_10.0
```

```
0%|          | 0/1000 [00:00<?, ?it/s]
```

```
training gauss_100.0
```

```
0%|          | 0/1000 [00:00<?, ?it/s]
```

```
[17]: # if you did everything correctly so far, this should output a nice figure you
      ↪ can use in your report
      plot_feature_mapping_comparison(outputs, y_test.reshape(size,size,3))
```




### 3 High Resolution Reconstruction

**High Resolution Reconstruction - Optimizer of your Choice - Various Input Mapping Strategies** Repeat the previous experiment, but at the higher resolution. The reason why we have you first experiment with the lower resolution since it is faster to train and debug. Additionally, you will see how the mapping strategies perform better or worse at the two different input resolutions.

```
[17]: size = 128
      train_data, test_data = get_image(size)
```

```
C:\Users\liang\AppData\Local\Temp\ipykernel_10408\3999599749.py:6:
DeprecationWarning: Starting with ImageIO v3 the behavior of this function will
switch to that of iio.v3.imread. To keep the current behavior (and make this
warning disappear) use `import imageio.v2 as imageio` or call
`imageio.v2.imread` directly.
  img = imageio.imread(image_url)[..., :3] / 255.
```



```
[20]: outputs = {}  
y_test = {}  
opt = "Adam"  
for k in tqdm(B_dict):  
    print("training", k)  
    X_train, y_train, X_test, y_test = get_input_features(B_dict, k)  
    outputs[k] = train_wrapper(k, size, opt)
```

```
0%|          | 0/5 [00:00<?, ?it/s]
```

```
training none
```

```
0%|          | 0/1000 [00:00<?, ?it/s]
```

```
training basic
```

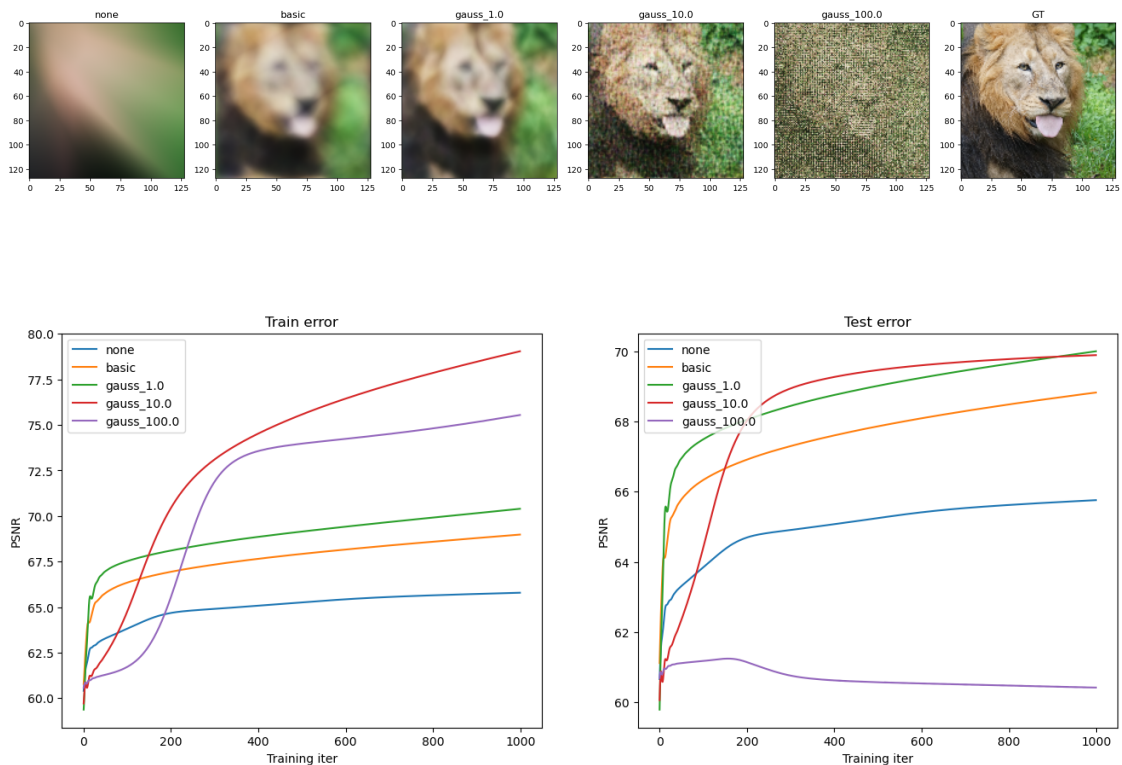


```

0%|          | 0/1000 [00:00<?, ?it/s]
training gauss_1.0
0%|          | 0/1000 [00:00<?, ?it/s]
training gauss_10.0
0%|          | 0/1000 [00:00<?, ?it/s]
training gauss_100.0
0%|          | 0/1000 [00:00<?, ?it/s]

```

```
[21]: plot_feature_mapping_comparison(outputs, y_test.reshape(size,size,3))
```



**High Resolution Reconstruction - Image of your Choice** When choosing an image select one that you think will give you interesting results or a better insight into the performance of different feature mappings and explain why in your report template.

```

[16]: size = 128
      # TODO pick an image and replace the url string

      train_data, test_data = get_image(size, image_url='https://bmild.github.io/
      ↪fourfeat/img/greece_orig.png')

```

```
C:\Users\liang\AppData\Local\Temp\ipykernel_22048\3999599749.py:6:
DeprecationWarning: Starting with ImageIO v3 the behavior of this function will
switch to that of iio.v3.imread. To keep the current behavior (and make this
warning disappear) use `import imageio.v2 as imageio` or call
`imageio.v2.imread` directly.
img = imageio.imread(image_url)[..., :3] / 255.
```



```
[17]: # get input features
      # TODO implement this by using the get_B_dict() and get_input_features() helper_
      ↪ functions
      mapping = 'gauss_10.0'
      B_dict = get_B_dict(128)
      X_train, y_train, X_test, y_test = get_input_features(B_dict, mapping)
```

```

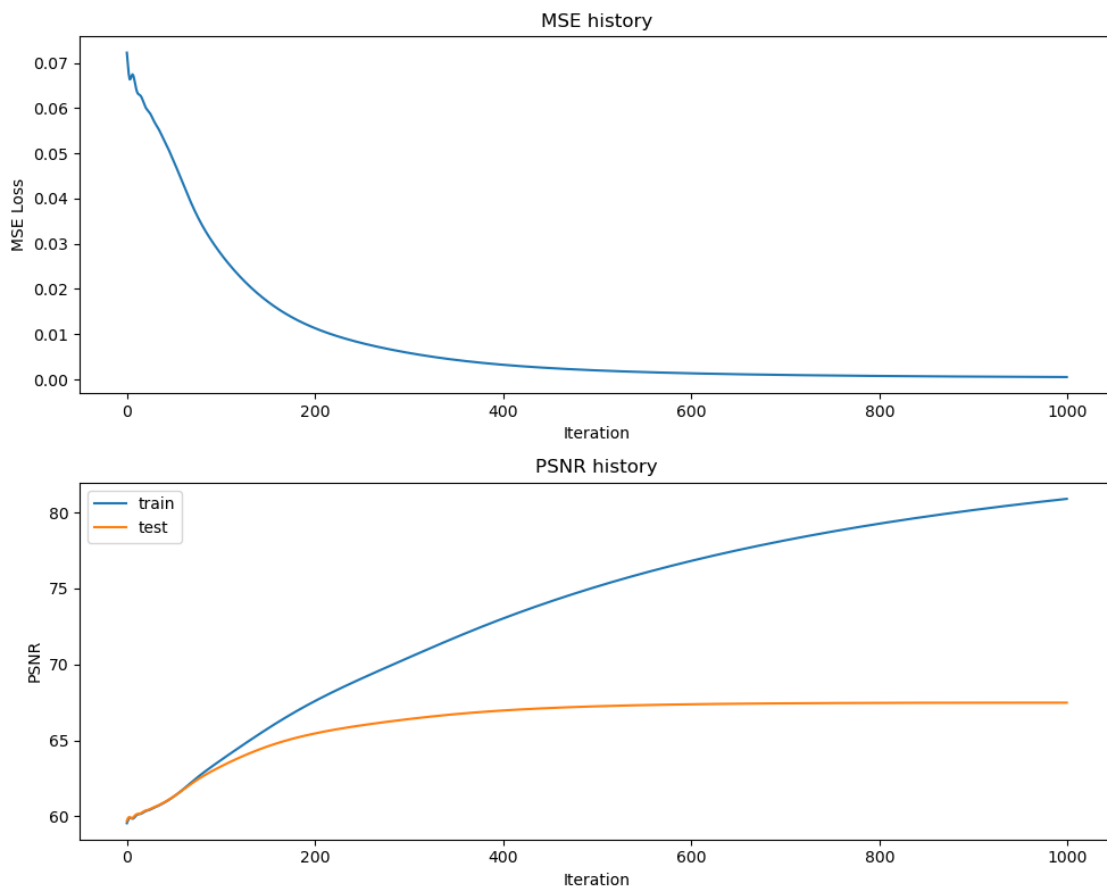
input_size = X_train.shape[1]
opt = "Adam"

# run NN experiment on input features
# TODO implement by using the NN_experiment() helper function
net, train_psnr, test_psnr, train_loss, predicted_images = \
    NN_experiment(X_train, y_train, X_test, y_test, input_size, num_layers, \
                  hidden_size, hidden_sizes, output_size, epochs, \
                  learning_rate, opt)

plot_training_curves(train_loss, train_psnr, test_psnr)
plot_reconstruction(net.forward(X_test), y_test)
plot_reconstruction_progress(predicted_images, y_test)

```

0%| | 0/1000 [00:00<?, ?it/s]



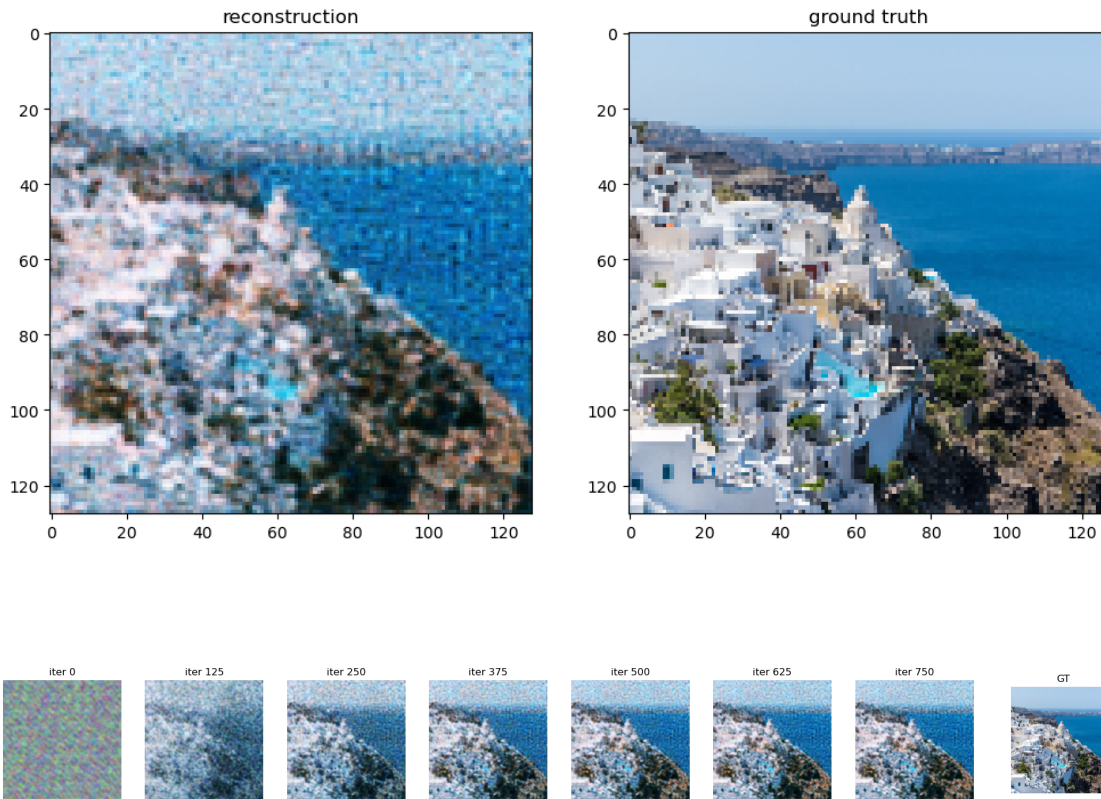
Final Test MSE 0.01162430334121582

Final Test psnr 67.47713426328009

C:\Users\liang\AppData\Local\Temp\ipykernel\_22048\3758130134.py:49:

MatplotlibDeprecationWarning: Auto-removal of overlapping axes is deprecated since 3.6 and will be removed two minor releases later; explicitly call `ax.remove()` as needed.

```
plt.subplot(1, N+1, N+1)
```



## 4 Reconstruction Process Video (Optional)

(For Fun!) Visualize the progress of training in a video

```
[ ]: # requires installing this additional dependency
!pip install imageio-ffmpeg
```

```
[ ]: # single video example
create_and_visualize_video({"gauss": {"pred_imgs": predicted_images}},  
    ↪filename="training_high_res_gauss.mp4")
```

```
[ ]: # multi video example
create_and_visualize_video(outputs, epochs=1000, size=32)
```