

## ▼ (Optional) Colab Setup

If you aren't using Colab, you can delete the following code cell. This is just to help students with mounting to Google Drive to access the other .py files and downloading the data, which is a little trickier on Colab than on your local machine using Jupyter.

```
1 # you will be prompted with a window asking to grant permissions
2 from google.colab import drive
3 drive.mount("/content/drive")

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

1 # fill in the path in your Google Drive in the string below. Note: do not escape slashes or spaces
2 import os
3 datadir = "/content/drive/MyDrive/CS444/assignment4"
4 if not os.path.exists(datadir):
5   !ln -s "/content/drive/MyDrive/path_to/assignment4" $datadir # TODO: Fill your Assignment 4 path
6 os.chdir(datadir)
7 !pwd

/content/drive/MyDrive/CS444/assignment4
```

## ▼ Generative Adversarial Networks

For this part of the assignment you implement two different types of generative adversarial networks. We will train the networks on a dataset of cat face images.

```
1 import torch
2 from torch.utils.data import DataLoader
3 from torchvision import transforms
4 from torchvision.datasets import ImageFolder
5 import matplotlib.pyplot as plt
6 import numpy as np
7
8 %matplotlib inline
9 plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
10 plt.rcParams['image.interpolation'] = 'nearest'
11 plt.rcParams['image.cmap'] = 'gray'
12
13 %load_ext autoreload
14 %autoreload 2
```

The autoreload extension is already loaded. To reload it, use:  
  %reload\_ext autoreload

```
1 from gan.train import train

1 device = torch.device("cuda:0" if torch.cuda.is_available() else "gpu")
```

## ▼ GAN loss functions

In this assignment you will implement two different types of GAN cost functions. You will first implement the loss from the [original GAN paper](#). You will also implement the loss from [LS-GAN](#).

### ▼ GAN loss

**TODO:** Implement the `discriminator_loss` and `generator_loss` functions in `gan/losses.py`.

The generator loss is given by:

$$\ell_G = -\mathbb{E}_{z \sim p(z)} [\log D(G(z))]$$

and the discriminator loss is:

$$\ell_D = -\mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] - \mathbb{E}_{z \sim p(z)} [\log(1 - D(G(z)))]$$

Note that these are negated from the equations presented earlier as we will be *minimizing* these losses.

**HINTS:** You should use the `torch.nn.functional.binary_cross_entropy_with_logits` function to compute the binary cross entropy loss since it is more numerically stable than using a softmax followed by BCE loss. The BCE loss is needed to compute the log probability of the true label given the logits output from the discriminator. Given a score  $s \in \mathbb{R}$  and a label  $y \in \{0, 1\}$ , the binary cross entropy loss is

$$bce(s, y) = -y * \log(s) - (1 - y) * \log(1 - s)$$

Instead of computing the expectation of  $\log D(G(z))$ ,  $\log D(x)$  and  $\log(1 - D(G(z)))$ , we will be averaging over elements of the minibatch, so make sure to combine the loss by averaging instead of summing.

```
1 from gan.losses import discriminator_loss, generator_loss
```

## ✓ Least Squares GAN loss

**TODO:** Implement the `ls_discriminator_loss` and `ls_generator_loss` functions in `gan/losses.py`.

We'll now look at [Least Squares GAN](#), a newer, more stable alternative to the original GAN loss function. For this part, all we have to do is change the loss function and retrain the model. We'll implement equation (9) in the paper, with the generator loss:

$$\ell_G = \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)) - 1)^2]$$

and the discriminator loss:

$$\ell_D = \frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}} [(D(x) - 1)^2] + \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)))^2]$$

**HINTS:** Instead of computing the expectation, we will be averaging over elements of the minibatch, so make sure to combine the loss by averaging instead of summing. When plugging in for  $D(x)$  and  $D(G(z))$  use the direct output from the discriminator (`scores_real` and `scores_fake`).

```
1 from gan.losses import ls_discriminator_loss, ls_generator_loss
```

## ✓ GAN model architecture

**TODO:** Implement the Discriminator and Generator networks in `gan/models.py`.

We recommend the following architectures which are inspired by [DCGAN](#):

### Discriminator:

- convolutional layer with `in_channels=3, out_channels=128, kernel=4, stride=2`
- convolutional layer with `in_channels=128, out_channels=256, kernel=4, stride=2`
- batch norm
- convolutional layer with `in_channels=256, out_channels=512, kernel=4, stride=2`
- batch norm
- convolutional layer with `in_channels=512, out_channels=1024, kernel=4, stride=2`
- batch norm
- convolutional layer with `in_channels=1024, out_channels=1, kernel=4, stride=1`

Use padding = 1 (not 0) for all the convolutional layers.

Instead of Relu we LeakyReLU throughout the discriminator (we use a negative slope value of 0.2). You can use simply use relu as well.

The output of your discriminator should be a single value score corresponding to each input sample. See `torch.nn.LeakyReLU`.

### Generator:

**Note:** In the generator, you will need to use transposed convolution (sometimes known as fractionally-strided convolution or deconvolution). This function is implemented in pytorch as `torch.nn.ConvTranspose2d`.

- transpose convolution with `in_channels=NOISE_DIM, out_channels=1024, kernel=4, stride=1`
- batch norm
- transpose convolution with `in_channels=1024, out_channels=512, kernel=4, stride=2`
- batch norm
- transpose convolution with `in_channels=512, out_channels=256, kernel=4, stride=2`
- batch norm
- transpose convolution with `in_channels=256, out_channels=128, kernel=4, stride=2`
- batch norm
- transpose convolution with `in_channels=128, out_channels=3, kernel=4, stride=2`

The output of the final layer of the generator network should have a `tanh` nonlinearity to output values between -1 and 1. The output should be a 3x64x64 tensor for each sample (equal dimensions to the images from the dataset).

```
1  from gan.models import Discriminator, Generator
```

## ▼ Data loading

The cat images we provide are RGB images with a resolution of 64x64. In order to prevent our discriminator from overfitting, we will need to perform some data augmentation.

**TODO:** Implement data augmentation by adding new transforms to the cell below. At the minimum, you should have a RandomCrop and a ColorJitter, but we encourage you to experiment with different augmentations to see how the performance of the GAN changes. See <https://pytorch.org/vision/stable/transforms.html>.

```
1  #!sh download_cat.sh

2
3
4
5
6
7
8
9
10
11
12
13
```

```
1  batch_size = 64
2  imsize = 64
3  cat_root = './cats'
4
5  cat_train = ImageFolder(root=cat_root, transform=transforms.Compose([
6      transforms.ToTensor(),
7
8      # Example use of RandomCrop:
9      transforms.Resize(int(1.15 * imsize)),
10     transforms.RandomCrop(imsize),
11 ])
12)
13  cat_loader_train = DataLoader(cat_train, batch_size=batch_size, drop_last=True)
```

## ▼ Visualize dataset

```
1 from gan.utils import show_images
2
3 try:
4     imgs = next(iter(cat_loader_train))[0].numpy().squeeze()
5 except:
6     imgs = cat_loader_train.__iter__().next()[0].numpy().squeeze()
7
8 show_images(imgs, color=True)
```



## ▼ Training

**TODO:** Fill in the training loop in `gan/train.py`.

```
1 NOISE_DIM = 100
2 NUM_EPOCHS = 50
3 learning_rate = 0.001
```

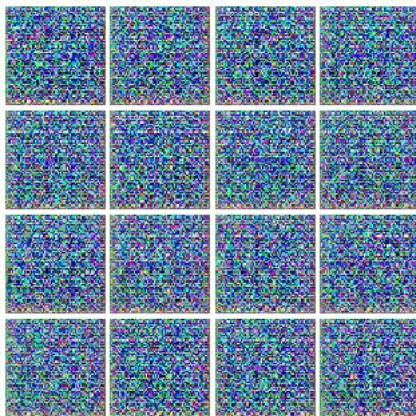
## ▼ Train GAN

```
1 D = Discriminator().to(device)
2 G = Generator(noise_dim=NOISE_DIM).to(device)

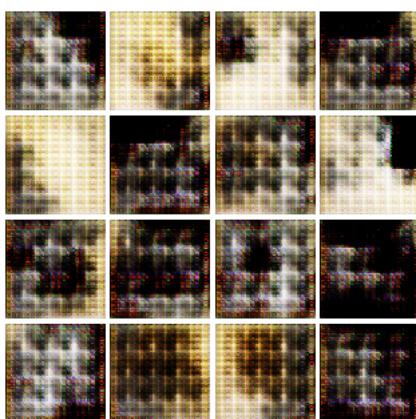
1 D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.999))
2 G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.999))

1 # original gan
2 train(D, G, D_optimizer, G_optimizer, discriminator_loss,
3       generator_loss, num_epochs=NUM_EPOCHS, show_every=250,
4       batch_size=batch_size, train_loader=cat_loader_train, device=device)
5
```

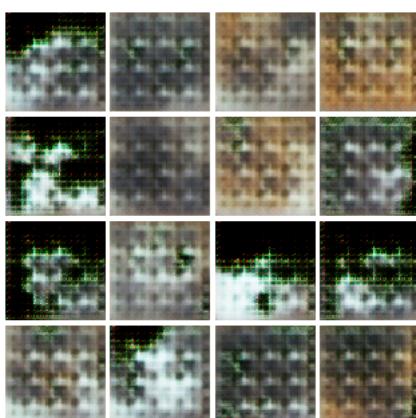
EPOCH: 1  
Iter: 0, D: 0.7079, G:6.138



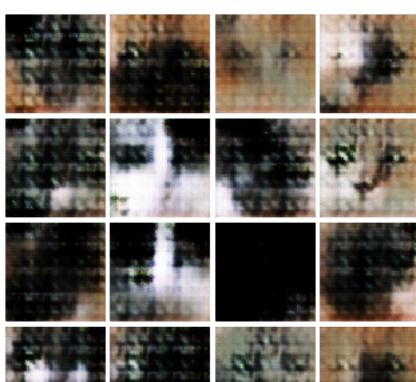
EPOCH: 2  
Iter: 250, D: 0.7495, G:1.464



EPOCH: 3  
Iter: 500, D: 0.5613, G:3.031



EPOCH: 4  
Iter: 750, D: 0.4069, G:1.824





EPOCH: 5  
Iter: 1000, D: 0.6298, G:1.781



EPOCH: 6  
Iter: 1250, D: 0.5896, G:2.079



EPOCH: 7  
Iter: 1500, D: 0.5268, G:3.207



EPOCH: 8  
Iter: 1750, D: 0.6613, G:0.8647





EPOCH: 9  
Iter: 2000, D: 0.5169, G:3.233



EPOCH: 10  
Iter: 2250, D: 0.3921, G:2.485

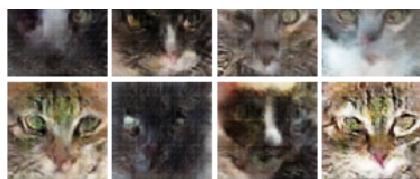


EPOCH: 11  
Iter: 2500, D: 0.763, G:6.523



EPOCH: 12  
Iter: 2750, D: 0.2897, G:3.461





EPOCH: 13  
Iter: 3000, D: 0.5748, G: 4.416



EPOCH: 14  
Iter: 3250, D: 0.3944, G: 2.004



EPOCH: 15  
Iter: 3500, D: 0.3888, G: 2.143



EPOCH: 16  
Iter: 3750, D: 0.9383, G: 1.232





EPOCH: 17

Iter: 4000, D: 0.3382, G:2.606



EPOCH: 18

Iter: 4250, D: 0.6438, G:1.417



EPOCH: 19

Iter: 4500, D: 0.2513, G:2.922



EPOCH: 20

Iter: 4750, D: 0.3874, G:3.204





EPOCH: 21  
Iter: 5000, D: 0.7917, G:2.148



EPOCH: 22  
Iter: 5250, D: 0.8108, G:0.8221



EPOCH: 23  
Iter: 5500, D: 0.1771, G:5.317



EPOCH: 24  
Iter: 5750, D: 0.2526, G:1.512





EPOCH: 25  
Iter: 6000, D: 0.1022, G: 6.074



EPOCH: 26  
Iter: 6250, D: 0.3238, G: 6.168



EPOCH: 27  
Iter: 6500, D: 0.3592, G: 7.43



EPOCH: 28  
Iter: 6750, D: 0.1865, G: 3.155

EPOCH: 0/50, D: 0.1000, G:5.400



EPOCH: 29

Iter: 7000, D: 0.09751, G:5.333



EPOCH: 30

Iter: 7250, D: 0.131, G:5.118



EPOCH: 31

Iter: 7500, D: 0.3644, G:8.199



EPOCH: 32

Iter: 7750, D: 0.08055, G:5.115



EPOCH: 33

Iter: 8000, D: 0.4334, G:2.318



EPOCH: 34

Iter: 8250, D: 0.956, G:8.119



EPOCH: 35

Iter: 8500, D: 0.1908, G:4.441





EPOCH: 36  
Iter: 8750, D: 0.2224, G: 2.082



EPOCH: 37  
Iter: 9000, D: 0.1188, G: 5.701



EPOCH: 38  
Iter: 9250, D: 0.9591, G: 10.47



EPOCH: 39  
Iter: 9500, D: 0.1058, G: 5.951





EPOCH: 40  
Iter: 9750, D: 0.04171, G:4.983



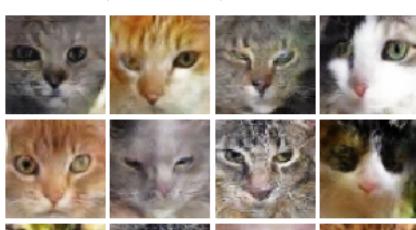
EPOCH: 41  
Iter: 10000, D: 0.06327, G:5.347

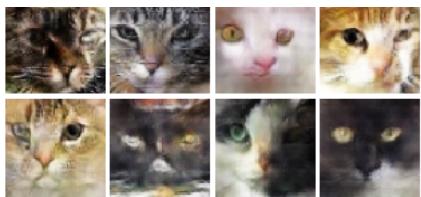


EPOCH: 42  
Iter: 10250, D: 0.1544, G:5.105



EPOCH: 43  
Iter: 10500, D: 0.4622, G:4.331





EPOCH: 44

Iter: 10750, D: 0.1333, G:4.409



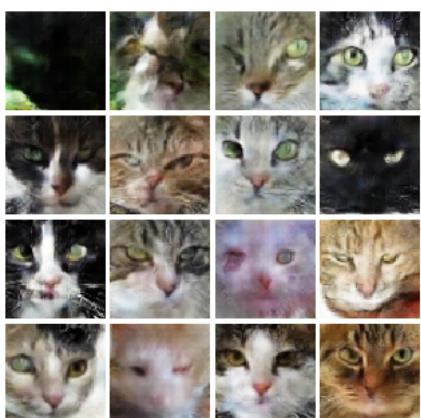
EPOCH: 45

Iter: 11000, D: 0.5665, G:12.09



EPOCH: 46

Iter: 11250, D: 1.263, G:16.85



EPOCH: 47

Iter: 11500, D: 0.09644, G:5.543





EPOCH: 48  
Iter: 11750, D: 0.2328, G:3.167



EPOCH: 49  
Iter: 12000, D: 0.06284, G:5.407



EPOCH: 50  
Iter: 12250, D: 0.8496, G:0.9974



## ✓ Train LS-GAN

```
1 D = Discriminator().to(device)
2 G = Generator(noise_dim=NOISE_DIM).to(device)

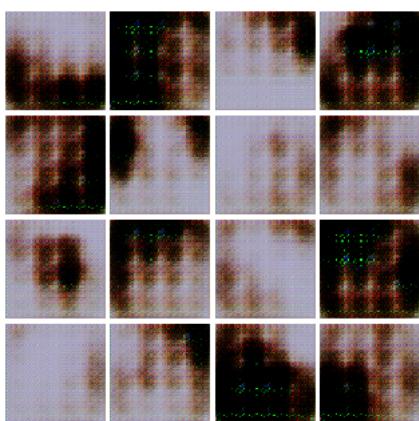
1 D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.999))
2 G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.999))

1 # ls-gan
2 train(D, G, D_optimizer, G_optimizer, ls_discriminator_loss,
3        ls_generator_loss, num_epochs=NUM_EPOCHS, show_every=250,
4        batch_size=batch_size, train_loader=cat_loader_train, device=device)
```

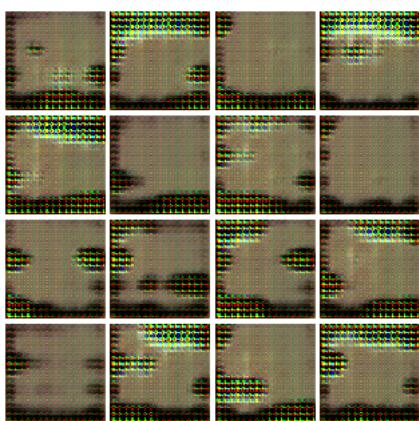
EPOCH: 1  
Iter: 0, D: 0.658, G:37.58



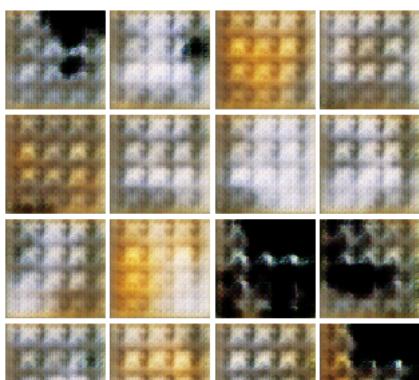
EPOCH: 2  
Iter: 250, D: 0.1837, G:0.5049



EPOCH: 3  
Iter: 500, D: 3.542, G:2.323



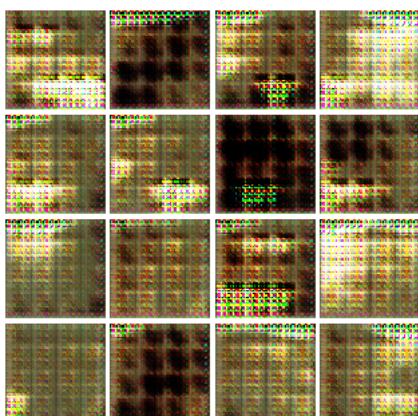
EPOCH: 4  
Iter: 750, D: 0.172, G:0.1387





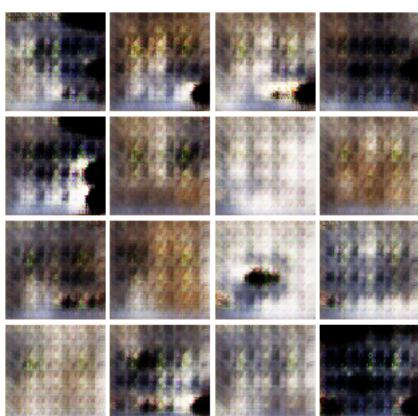
EPOCH: 5

Iter: 1000, D: 0.1274, G:0.2906



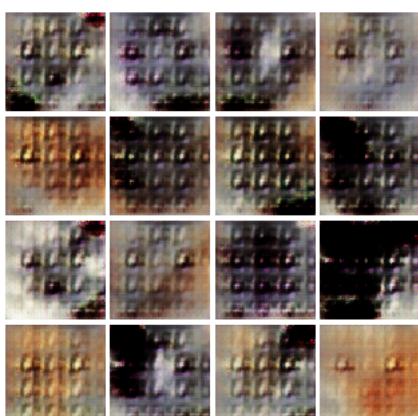
EPOCH: 6

Iter: 1250, D: 0.2233, G:0.2264



EPOCH: 7

Iter: 1500, D: 0.154, G:0.4346



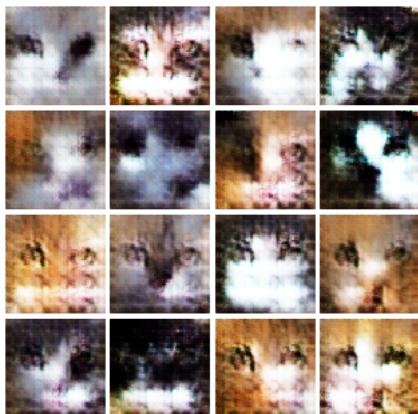
EPOCH: 8

Iter: 1750, D: 0.2533, G:0.2115

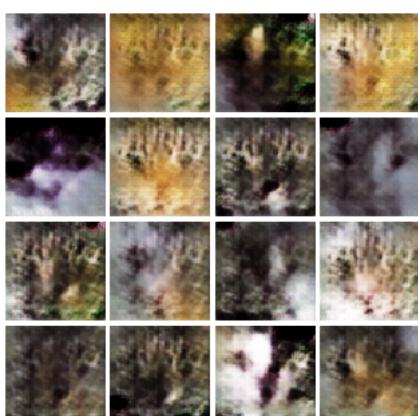




EPOCH: 9  
Iter: 2000, D: 0.1987, G:0.1856



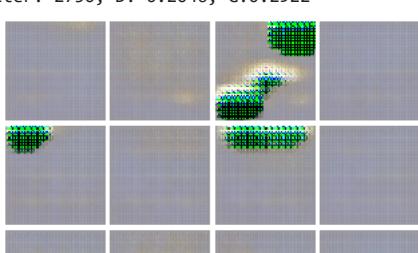
EPOCH: 10  
Iter: 2250, D: 0.3094, G:0.5547

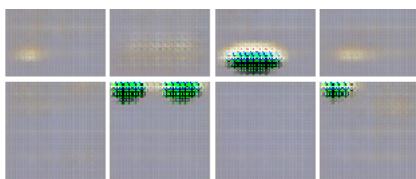


EPOCH: 11  
Iter: 2500, D: 0.2129, G:0.2157

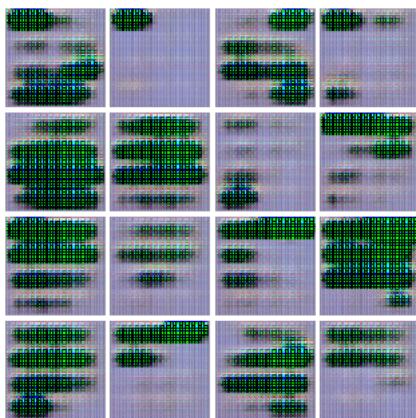


EPOCH: 12  
Iter: 2750, D: 0.2046, G:0.2922

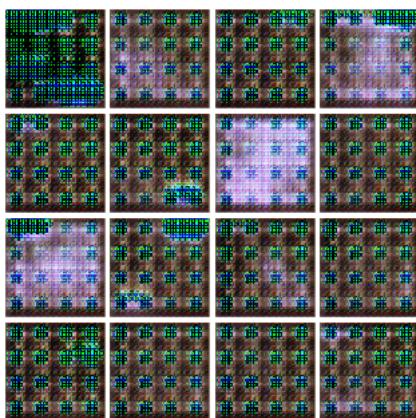




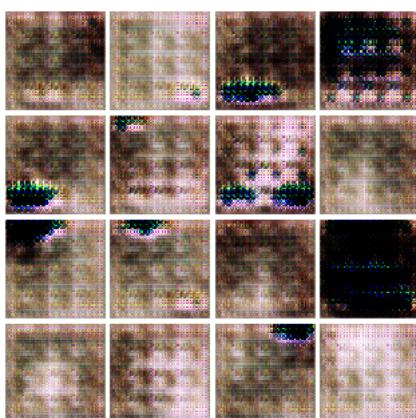
EPOCH: 13  
Iter: 3000, D: 0.1734, G: 0.1995



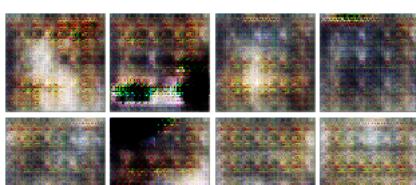
EPOCH: 14  
Iter: 3250, D: 0.1306, G: 0.2471

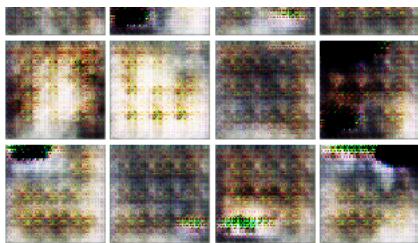


EPOCH: 15  
Iter: 3500, D: 0.2343, G: 0.1571



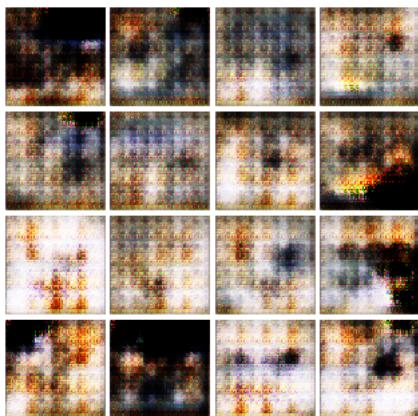
EPOCH: 16  
Iter: 3750, D: 0.224, G: 0.1262





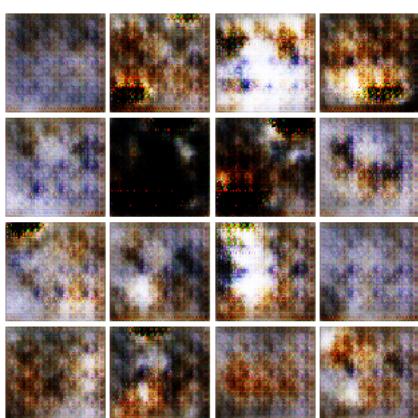
EPOCH: 17

Iter: 4000, D: 0.2781, G: 0.1489



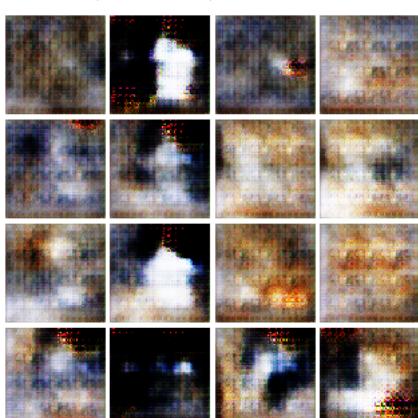
EPOCH: 18

Iter: 4250, D: 0.2575, G: 0.1378



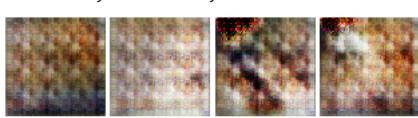
EPOCH: 19

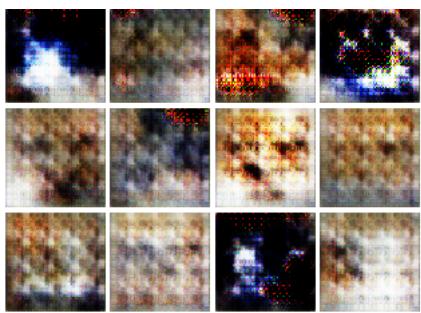
Iter: 4500, D: 0.2855, G: 0.1724



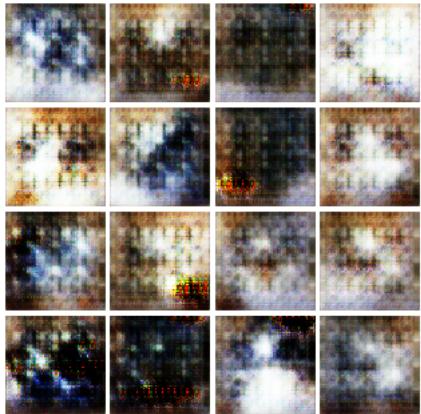
EPOCH: 20

Iter: 4750, D: 0.2466, G: 0.143

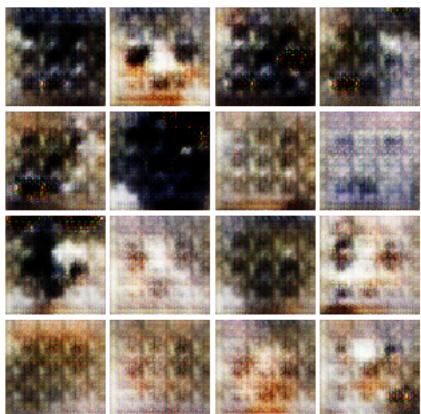




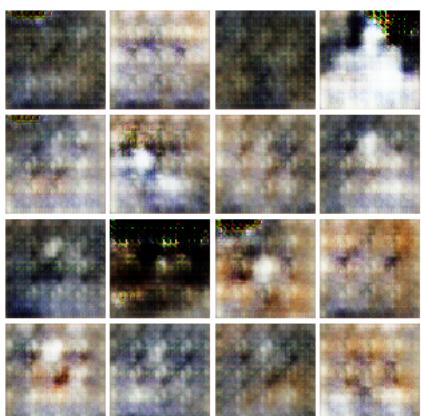
EPOCH: 21  
Iter: 5000, D: 0.3408, G: 0.2401



EPOCH: 22  
Iter: 5250, D: 0.3061, G: 0.2399

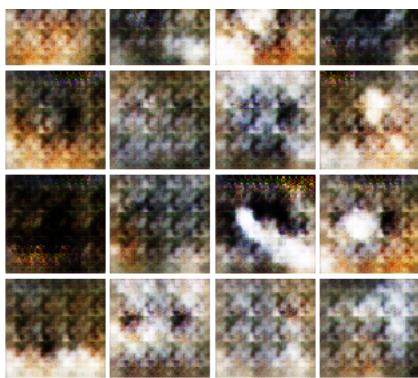


EPOCH: 23  
Iter: 5500, D: 0.2643, G: 0.1174

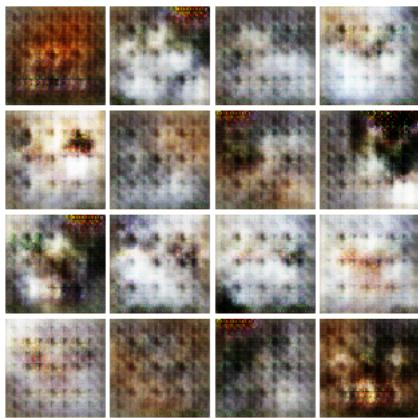


EPOCH: 24  
Iter: 5750, D: 0.2575, G: 0.1795

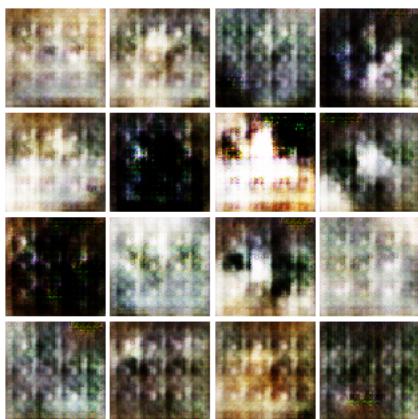




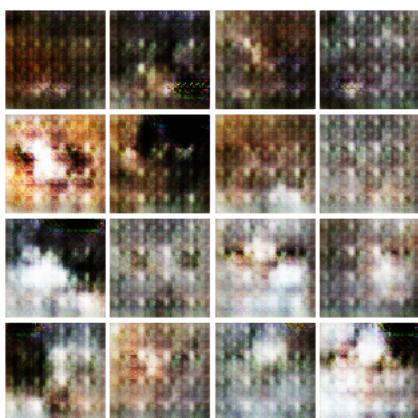
EPOCH: 25  
Iter: 6000, D: 0.281, G: 0.1825



EPOCH: 26  
Iter: 6250, D: 0.4961, G: 0.3849

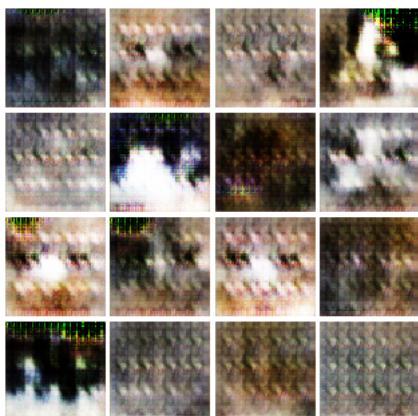


EPOCH: 27  
Iter: 6500, D: 0.5924, G: 0.2854



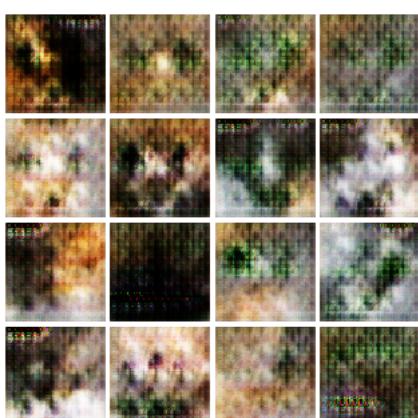
EPOCH: 28  
Iter: 6750, D: 0.2474, G: 0.1602

EPOCH: 2700, D: 0.2474, G:0.1602



EPOCH: 29

Iter: 7000, D: 0.2561, G:0.1579



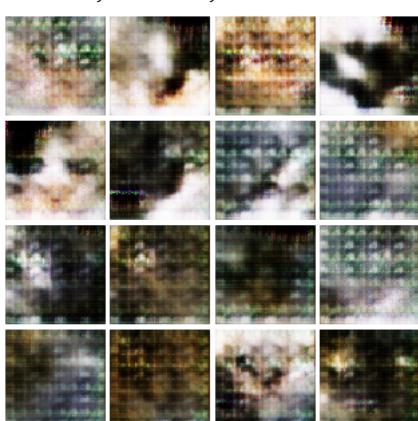
EPOCH: 30

Iter: 7250, D: 0.3258, G:0.2041



EPOCH: 31

Iter: 7500, D: 0.2492, G:0.2434



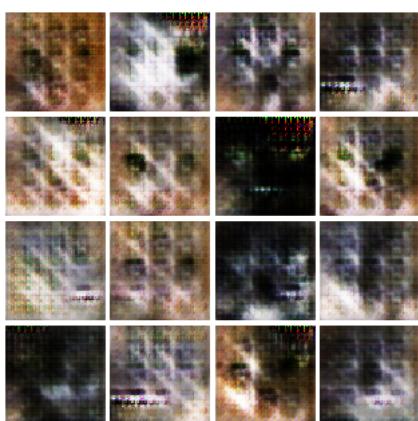
EPOCH: 32

Iter: 7750, D: 0.2388, G:0.1699



EPOCH: 33

Iter: 8000, D: 0.2549, G:0.2432



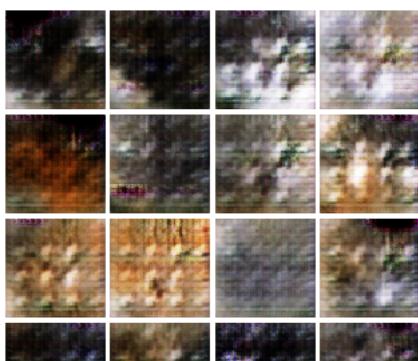
EPOCH: 34

Iter: 8250, D: 0.2484, G:0.1507



EPOCH: 35

Iter: 8500, D: 0.2662, G:0.17





EPOCH: 36  
Iter: 8750, D: 0.2502, G: 0.1812



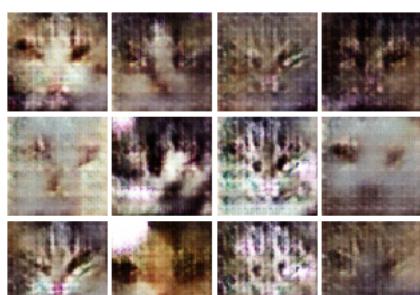
EPOCH: 37  
Iter: 9000, D: 0.2415, G: 0.1559



EPOCH: 38  
Iter: 9250, D: 0.2814, G: 0.2005



EPOCH: 39  
Iter: 9500, D: 0.2627, G: 0.2152





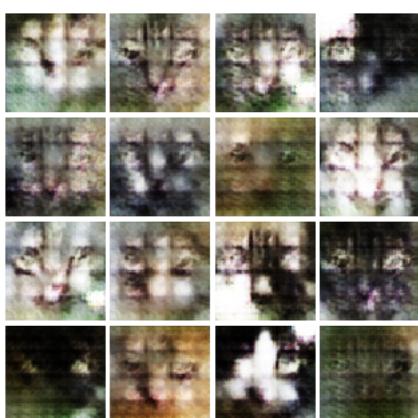
EPOCH: 40  
Iter: 9750, D: 0.2402, G:0.2164



EPOCH: 41  
Iter: 10000, D: 0.2854, G:0.1277

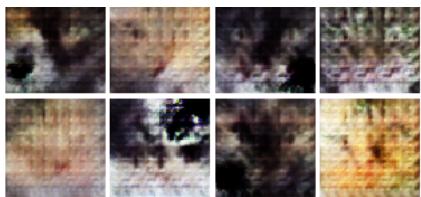


EPOCH: 42  
Iter: 10250, D: 0.2473, G:0.2101



EPOCH: 43  
Iter: 10500, D: 0.368, G:0.1819





EPOCH: 44

Iter: 10750, D: 0.2888, G:0.2008



EPOCH: 45

Iter: 11000, D: 0.2471, G:0.248



EPOCH: 46

Iter: 11250, D: 0.2249, G:0.2138

